

```
In [27]: import numpy as np
```

```
In [28]: import pandas as pd
```

```
In [29]: import matplotlib.pyplot as plt
```

```
In [30]: import seaborn as sns
```

```
In [31]: %matplotlib inline
```

```
In [32]: import warnings  
warnings.filterwarnings('ignore')
```

```
In [33]: df=pd.read_csv(r"D:\Data Science with AI\Data Science With AI\3rd, 4th -august-
```

```
In [34]: df.shape
```

```
Out[34]: (32561, 15)
```

```
In [35]: df.head()
```

```
Out[35]:   age  workclass  fnlwgt  education  education.num  marital.status  occupation  relati  
  
          0    90        ?    77053    HS-grad           9      Widowed          ?  
          1    82    Private  132870    HS-grad           9      Widowed  Exec-  
          2    66        ?   186061  Some-  
                           college           10      Widowed          ?      Unr  
          3    54    Private  140359    7th-8th           4      Divorced  Machine-  
                           op-inspt           10      Separated  Prof-  
                           specialty          Ow
```

```
In [36]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   age               32561 non-null   int64  
 1   workclass         32561 non-null   object  
 2   fnlwgt            32561 non-null   int64  
 3   education         32561 non-null   object  
 4   education.num    32561 non-null   int64  
 5   marital.status   32561 non-null   object  
 6   occupation        32561 non-null   object  
 7   relationship      32561 non-null   object  
 8   race               32561 non-null   object  
 9   sex                32561 non-null   object  
 10  capital.gain     32561 non-null   int64  
 11  capital.loss     32561 non-null   int64  
 12  hours.per.week   32561 non-null   int64  
 13  native.country   32561 non-null   object  
 14  income             32561 non-null   object  
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
In [37]: df[df=='?']=np.nan
```

```
In [38]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   age               32561 non-null   int64  
 1   workclass         30725 non-null   object  
 2   fnlwgt            32561 non-null   int64  
 3   education         32561 non-null   object  
 4   education.num    32561 non-null   int64  
 5   marital.status   32561 non-null   object  
 6   occupation        30718 non-null   object  
 7   relationship      32561 non-null   object  
 8   race               32561 non-null   object  
 9   sex                32561 non-null   object  
 10  capital.gain     32561 non-null   int64  
 11  capital.loss     32561 non-null   int64  
 12  hours.per.week   32561 non-null   int64  
 13  native.country   31978 non-null   object  
 14  income             32561 non-null   object  
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
In [39]: for col in ['workclass','occupation','native.country']:
    df[col].fillna(df[col].mode()[0],inplace=True)
```

```
In [40]: df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
0	90	Private	77053	HS-grad		9	Widowed	Prof-specialty
1	82	Private	132870	HS-grad		9	Widowed	Exec-managerial
2	66	Private	186061	Some-college		10	Widowed	Prof-specialty
3	54	Private	140359	7th-8th		4	Divorced	Machine-op-inspct
4	41	Private	264663	Some-college		10	Separated	Prof-specialty

◀ ▶

In [41]: `df.isnull().sum()`

```
Out[41]: age          0
          workclass     0
          fnlwgt        0
          education      0
          education.num  0
          marital.status 0
          occupation      0
          relationship    0
          race           0
          sex            0
          capital.gain   0
          capital.loss   0
          hours.per.week 0
          native.country  0
          income          0
          dtype: int64
```

In [42]: `x=df.drop(['income'],axis=1)`
`y=df['income']`

In [43]: `x`

Out[43]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	...
0	90	Private	77053	HS-grad		9	Widowed	Prof-specialty
1	82	Private	132870	HS-grad		9	Widowed	Exec-managerial
2	66	Private	186061	Some-college		10	Widowed	Prof-specialty
3	54	Private	140359	7th-8th		4	Divorced	Machine-op-inspct
4	41	Private	264663	Some-college		10	Separated	Prof-specialty
...
32556	22	Private	310152	Some-college		10	Never-married	Protective-serv
32557	27	Private	257302	Assoc-acdm		12	Married-civ-spouse	Tech-support
32558	40	Private	154374	HS-grad		9	Married-civ-spouse	Machine-op-inspct
32559	58	Private	151910	HS-grad		9	Widowed	Adm-clerical
32560	22	Private	201490	HS-grad		9	Never-married	Adm-clerical

32561 rows × 14 columns

In [44]: y

Out[44]:

0	<=50K
1	<=50K
2	<=50K
3	<=50K
4	<=50K
...	
32556	<=50K
32557	<=50K
32558	>50K
32559	<=50K
32560	<=50K

Name: income, Length: 32561, dtype: object

In [45]: x.head()

Out[45]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
0	90	Private	77053	HS-grad		9	Widowed	Prof-specialty
1	82	Private	132870	HS-grad		9	Widowed	Exec-managerial
2	66	Private	186061	Some-college		10	Widowed	Prof-specialty
3	54	Private	140359	7th-8th		4	Divorced	Machine-op-inspct
4	41	Private	264663	Some-college		10	Separated	Prof-specialty

In [46]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
```

In [47]:

```
from sklearn import preprocessing
categorical=['workclass','education','marital.status','occupation','relationship']
for feature in categorical:
    le=preprocessing.LabelEncoder()
    x_train[feature]=le.fit_transform(x_train[feature])
    x_test[feature]=le.transform(x_test[feature])
```

In [48]:

```
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
x_train=pd.DataFrame(scaler.fit_transform(x_train),columns=x.columns)
x_test=pd.DataFrame(scaler.transform(x_test),columns=x.columns)
```

In [49]:

```
x_train.head()
```

Out[49]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation
0	0.101484	2.600478	-1.494279	-0.332263	1.133894	-0.402341	-0.78223
1	0.028248	-1.884720	0.438778	0.184396	-0.423425	-0.402341	-0.02669
2	0.247956	-0.090641	0.045292	1.217715	-0.034095	0.926666	-0.78223
3	-0.850587	-1.884720	0.793152	0.184396	-0.423425	0.926666	-0.53038
4	-0.044989	-2.781760	-0.853275	0.442726	1.523223	-0.402341	-0.78223

Logistic Regression model with all features

In [50]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
logreg=LogisticRegression()
logreg.fit(x_train,y_train)
```

```
y_pred=logreg.predict(x_test)
print('Logistic Regression accuracy score with all the features:{0:0.4f}'.format
```

Logistic Regression accuracy score with all the features:0.8218

Logistic regression with PCA

```
In [25]: from sklearn.decomposition import PCA
pca=PCA()
x_train=pca.fit_transform(x_train)
pca.explained_variance_
```

```
Out[25]: array([0.14757168, 0.10182915, 0.08147199, 0.07880174, 0.07463545,
   0.07274281, 0.07009602, 0.06750902, 0.0647268 , 0.06131155,
   0.06084207, 0.04839584, 0.04265038, 0.02741548])
```

```
In [26]: x.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	
0	90	Private	77053	HS-grad		9	Widowed	Prof-specialty	
1	82	Private	132870	HS-grad		9	Widowed	Exec-managerial	
2	66	Private	186061	Some-college		10	Widowed	Prof-specialty	Unr
3	54	Private	140359	7th-8th		4	Divorced	Machine-op-inspct	Unr
4	41	Private	264663	Some-college		10	Separated	Prof-specialty	Ow

Logistic Regression with first 13 features

```
In [52]: x=df.drop(['income','native.country'],axis=1)
y=df['income']

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
categorical=['workclass','education','marital.status','occupation','relationship']
for feature in categorical:
    le=preprocessing.LabelEncoder()
    x_train[feature]=le.fit_transform(x_train[feature])
    x_test[feature]=le.transform(x_test[feature])
x_train=pd.DataFrame(scaler.fit_transform(x_train),columns=x.columns)
x_test=pd.DataFrame(scaler.transform(x_test),columns=x.columns)
logreg=LogisticRegression()
logreg.fit(x_train,y_train)
y_pred=logreg.predict(x_test)
print('Logistic Regression accuracy score with the first 13 features:{0:0.4f}'.f
```

Logistic Regression accuracy score with the first 13 features:0.8213

Logistic Regression with first 12 features

```
In [54]: x=df.drop(['income','native.country','hours.per.week'],axis=1)
y=df['income']

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
categorical=['workclass','education','marital.status','occupation','relationship'
for feature in categorical:
    le=preprocessing.LabelEncoder()
    x_train[feature]=le.fit_transform(x_train[feature])
    x_test[feature]=le.transform(x_test[feature])

x_train=pd.DataFrame(scaler.fit_transform(x_train),columns=x.columns)
x_test=pd.DataFrame(scaler.transform(x_test),columns=x.columns)

logreg=LogisticRegression()
logreg.fit(x_train,y_train)
y_pred=logreg.predict(x_test)
print('Logistic Regression accuracy score with the first 12 features:{0:0.4f}'.f
```

Logistic Regression accuracy score with the first 12 features:0.8227

Logistic Regression with first 11 features

```
In [56]: x=df.drop(['income','native.country','hours.per.week','capital.loss'],axis=1)
y=df['income']

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)

categorical=['workclass','education','marital.status','occupation','relationship'
for feature in categorical:
    le=preprocessing.LabelEncoder()
    x_train[feature]=le.fit_transform(x_train[feature])
    x_test[feature]=le.transform(x_test[feature])

x_train=pd.DataFrame(scaler.fit_transform(x_train),columns=x.columns)
x_test=pd.DataFrame(scaler.transform(x_test),columns=x.columns)

logreg=LogisticRegression()
logreg.fit(x_train,y_train)
y_pred=logreg.predict(x_test)

print('Logistic Regression accuracy score with the first 11 features:{0:0.4f}'.f
```

Logistic Regression accuracy score with the first 11 features:0.8186

select right number of dimensions

```
In [57]: x=df.drop(['income'],axis=1)
y=df['income']

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
```

```
categorical=['workclass','education','marital.status','occupation','relationship'
for feature in categorical:
    le=preprocessing.LabelEncoder()
    x_train[feature]=le.fit_transform(x_train[feature])
    x_test[feature]=le.transform(x_test[feature])

x_train=pd.DataFrame(scaler.fit_transform(x_train),columns=x.columns)

pca=PCA()
pca.fit(x_train)
cumsum=np.cumsum(pca.explained_variance_ratio_)
dim=np.argmax(cumsum>=0.90)+1
print('The number of dimensions required to preserve 90% of variance is',dim)
```

The number of dimensions required to preserve 90% of variance is 12

In []: